




INTRODUCTION TO DATA SCIENCE

WITH PYTHON



Welcome!

Please find a
seat and get
settled

Start your laptop
and check
internet

We will start at
9.00am!


Welcome!

Presenter introduction


- Lindsay Forestell
- Stuart Hemerling

Acknowledgement

- [Indigenous Peoples and Lands \(cirnac.gc.ca\)](http://cirnac.gc.ca)
- [The First Nations Principles of OCAP® - The First Nations Information Governance Centre \(fnigc.ca\)](http://fnigc.ca)
- [Native-Land.ca | Our home on native land](http://Native-Land.ca)



The Learning Environment



1 Content is new, developed by the DSP team

2 Geared to the practice of data science

3 Course is interactive – no textbook – you will be coding and saving your work for future reference

4 We will aim to start/end sessions on time

Day 1

- 9:00-9:30 - Course Introduction (20 min) (Stuart)
- 9:30-10:30 - [Getting Up and Running](#) (60 min) (Lindsay)
- 10:30-10:45 BREAK ☕
- 10:45-Noon [Working with Code](#) (90 min) (Stuart)
- Noon-1:00 LUNCH 🍌
- 1:00-2:30 [Core Data Structure Concepts](#) (90 min) (Lindsay)
- 2:30-2:45 BREAK 🎧
- 2:45-4:30 [Getting Data](#) (90 min) (Lindsay)

Day 2

- 9:00-10:30 - [Data Cleaning](#) (90 min) (Stuart)
- 10:30-10:45 BREAK ☕
- 10:45-Noon [Exploring Data with Pandas](#) (90 min) (Stuart)
- Noon-1:00 LUNCH 🍌
- 1:00-2:30 [Graphical Depictions of Data](#) (90 min) (Lindsay)
- 2:30-2:45 BREAK 🎧
- 2:45-4:30 Bonus Content: TBD (90 min) (TBD - Learners vote on topic end of Day 1!)

This is a general guide of what will be covered - some sections may start or end earlier or later than outlined!



Sticky Notes

1

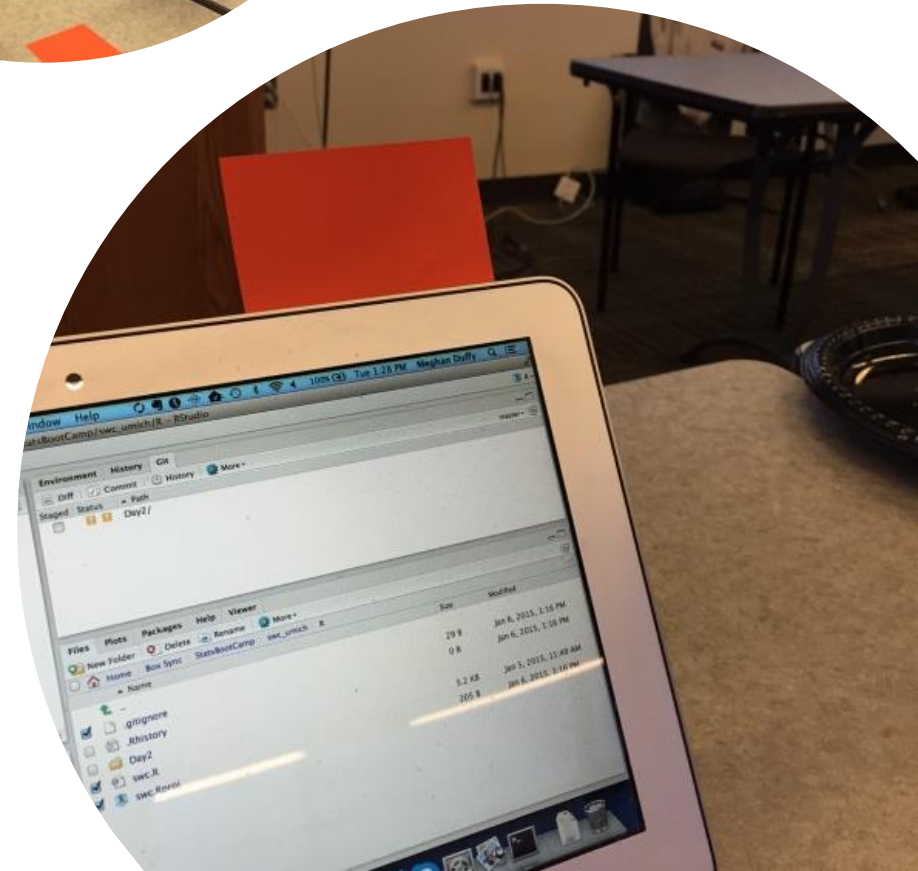
- As Status Flags
 - Green is "*I'm done!*"
 - Red is "*I need help!*"
 - No flag means "*Working on it*"

2


3

4

- Please put your name on both so we can get to know each other properly!



Collaborative Note Taking: Etherpad

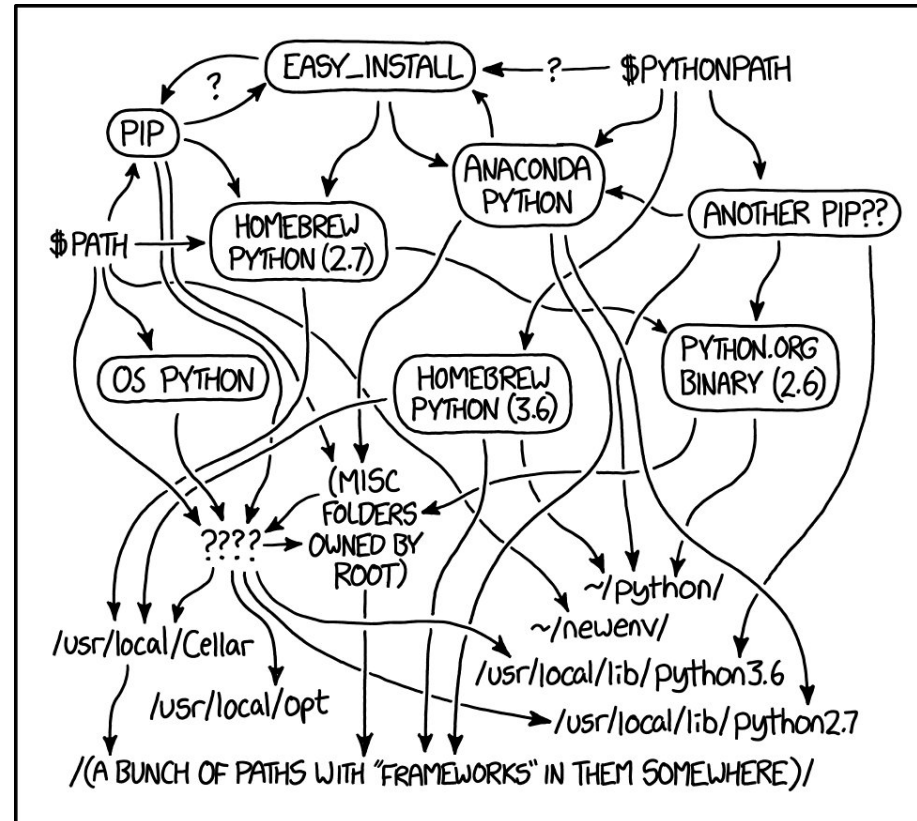



- <https://etherpad.wikimedia.org/p/bcgov-intro-python-jan23>
- Where we can share challenges and solutions throughout course
- Let's start by going there!

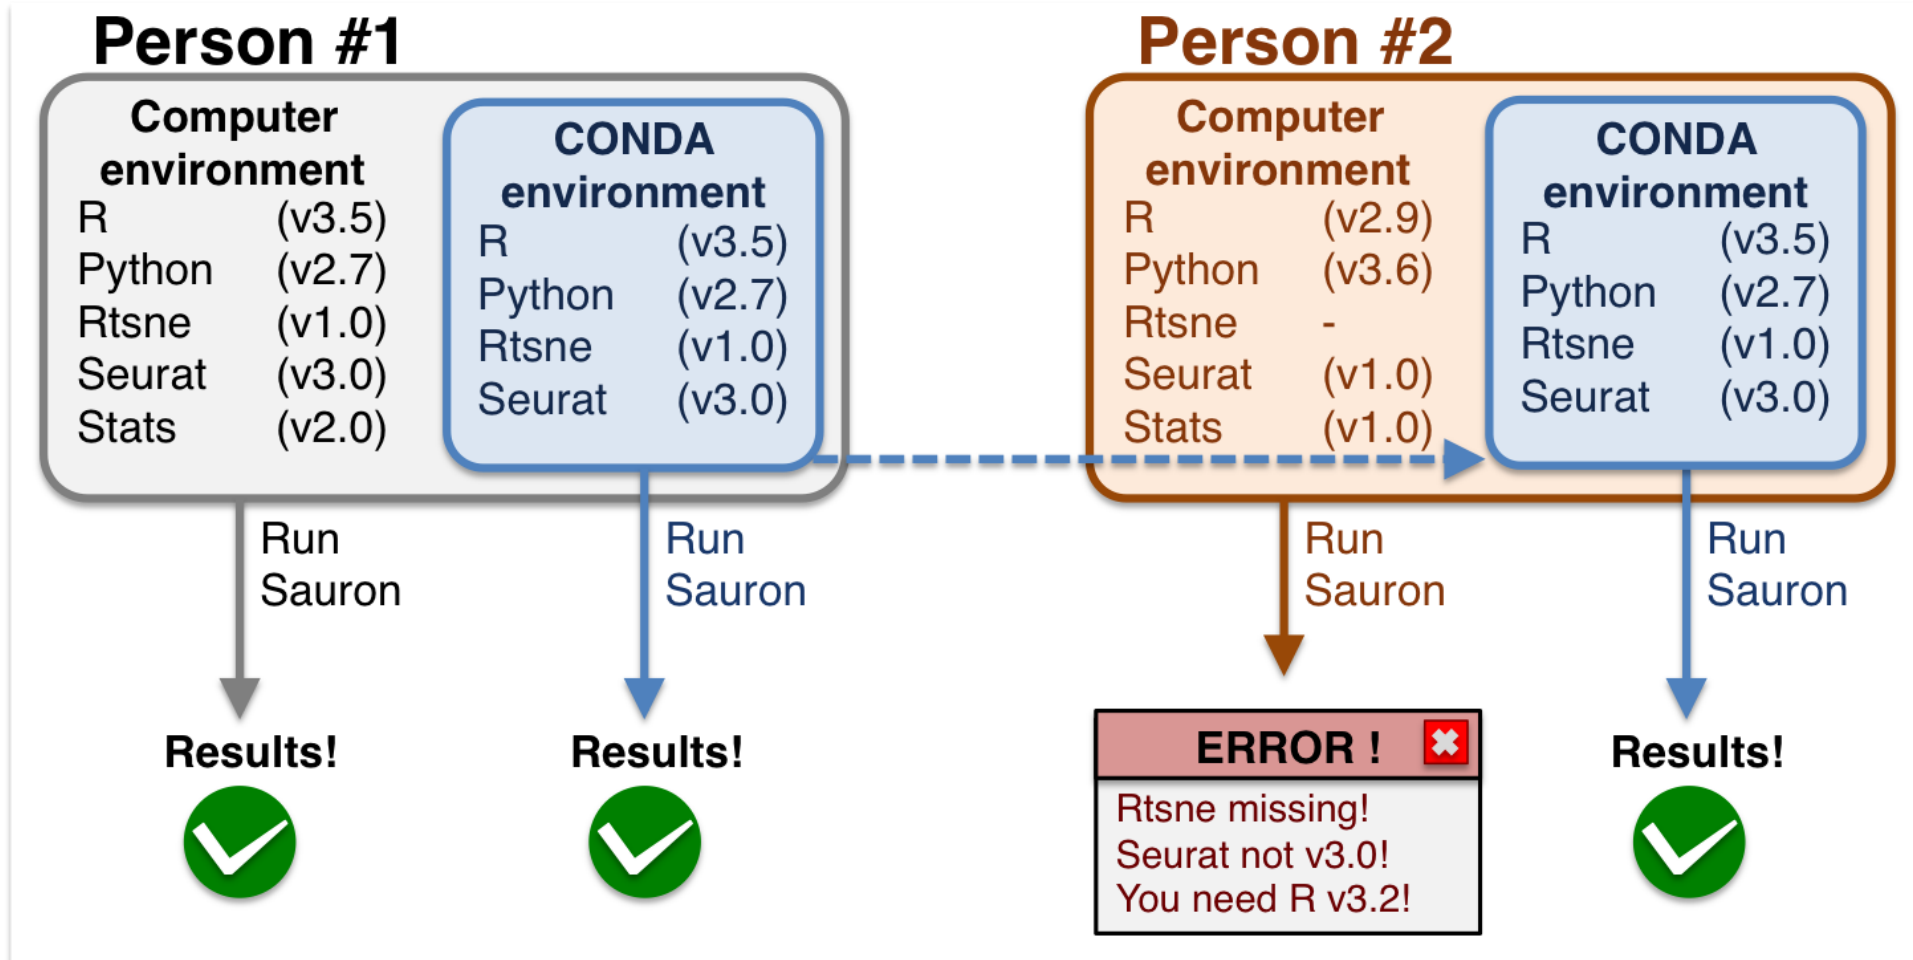
GETTING UP AND RUNNING



Why use conda?



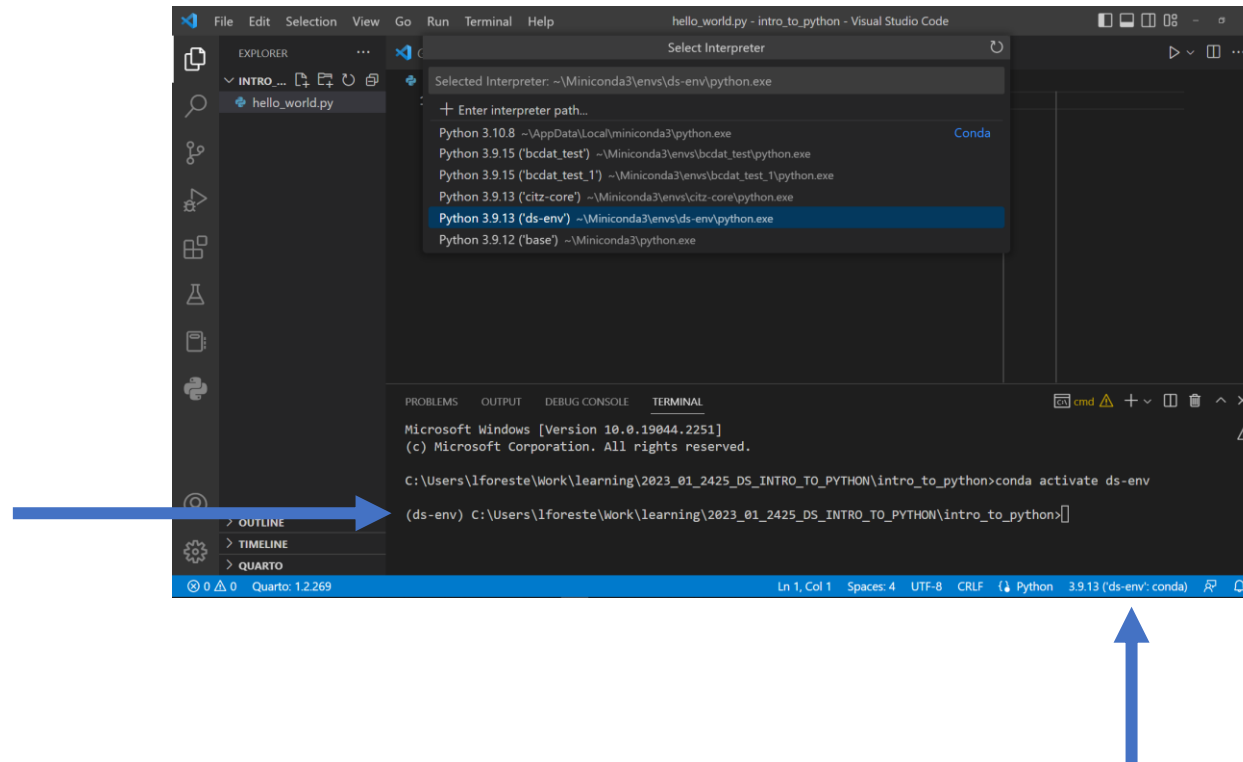
Why use conda?




Remember to double check:

Does my command line start with the right environment name?

Is VS Code using the right environment?



Different Environments for Different Purposes



Jupyter Lab

- Barebones IDE
- Useful for exploratory and documented analysis done in **notebooks**
- .ipynb files
- Displays plots inline, between **cells** of code

VS Code

- More wholistic IDE (integrated developing environment)
- Useful for running, testing, debugging **scripts**
- .py files
- Displays plots externally
- No notion of code cells

Hello, world!

1

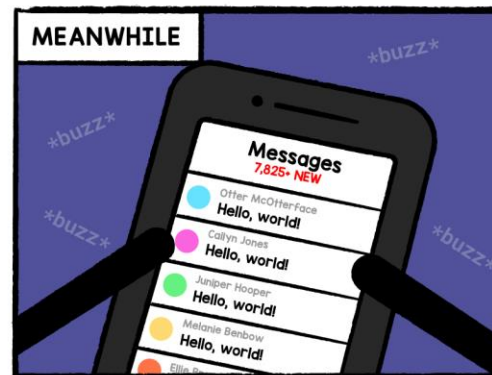

2

3

4

OTTER THIS WORLD


REIKACANJA @REIKACANJA



Challenge 1

In the **hello_world.py** file, run the second line of code.


Add a third line of code that prints your favourite TV show, and print this line.



Challenge 2


Try clicking the play button again.

What happens here? Can you explain why?




Challenge 3

Rename the notebook to **hello_world.ipynb**



Challenge 4


Print 'Hello World' inside the Jupyter Notebook.



WORKING WITH CODE



Organizing Matters




SEE, I'VE GOT A REALLY GOOD SYSTEM:
IF I WANT TO SEND A YOUTUBE VIDEO
TO SOMEONE, I GO TO FILE → SAVE, THEN
IMPORT THE SAVED PAGE INTO WORD. THEN
I GO TO "SHARE THIS DOCUMENT" AND
UNDER "RECIPIENT" I PUT THE EMAIL
OF THIS VIDEO EXTRACTION SERVICE...




I'LL OFTEN ENCOURAGE RELATIVES TO TRY TO SOLVE
COMPUTER PROBLEMS THEMSELVES BY TRIAL AND ERROR.
HOWEVER, I'VE LEARNED AN IMPORTANT LESSON: IF THEY
SAY THEY'VE SOLVED THEIR PROBLEM, *NEVER* ASK HOW.

Organizing Principles



1. Treat data as read only
2. Clean and analyze your data with code
3. Treat generated output as disposable
4. Organize your files


File Organization "Template"



/MarmotsBC	top level directory, contains project name, without spaces
— /docs	supporting text and word processed documents that are meant for humans to read
— /data	the original (meta)data that acts as an input and will be treated as read-only
— /src	your scripts, may further be divided into /functions, /test, etc.
— /bin	programs brought in from elsewhere
— /output	data written by your code to this location, could be input for another process
— /analysis	the charts, stats, figures and the like generated from your scripts


Challenge 1

What are three activities associated with a project that are made more difficult and put at risk if project files are not kept well organized? Who is likely to be affected?



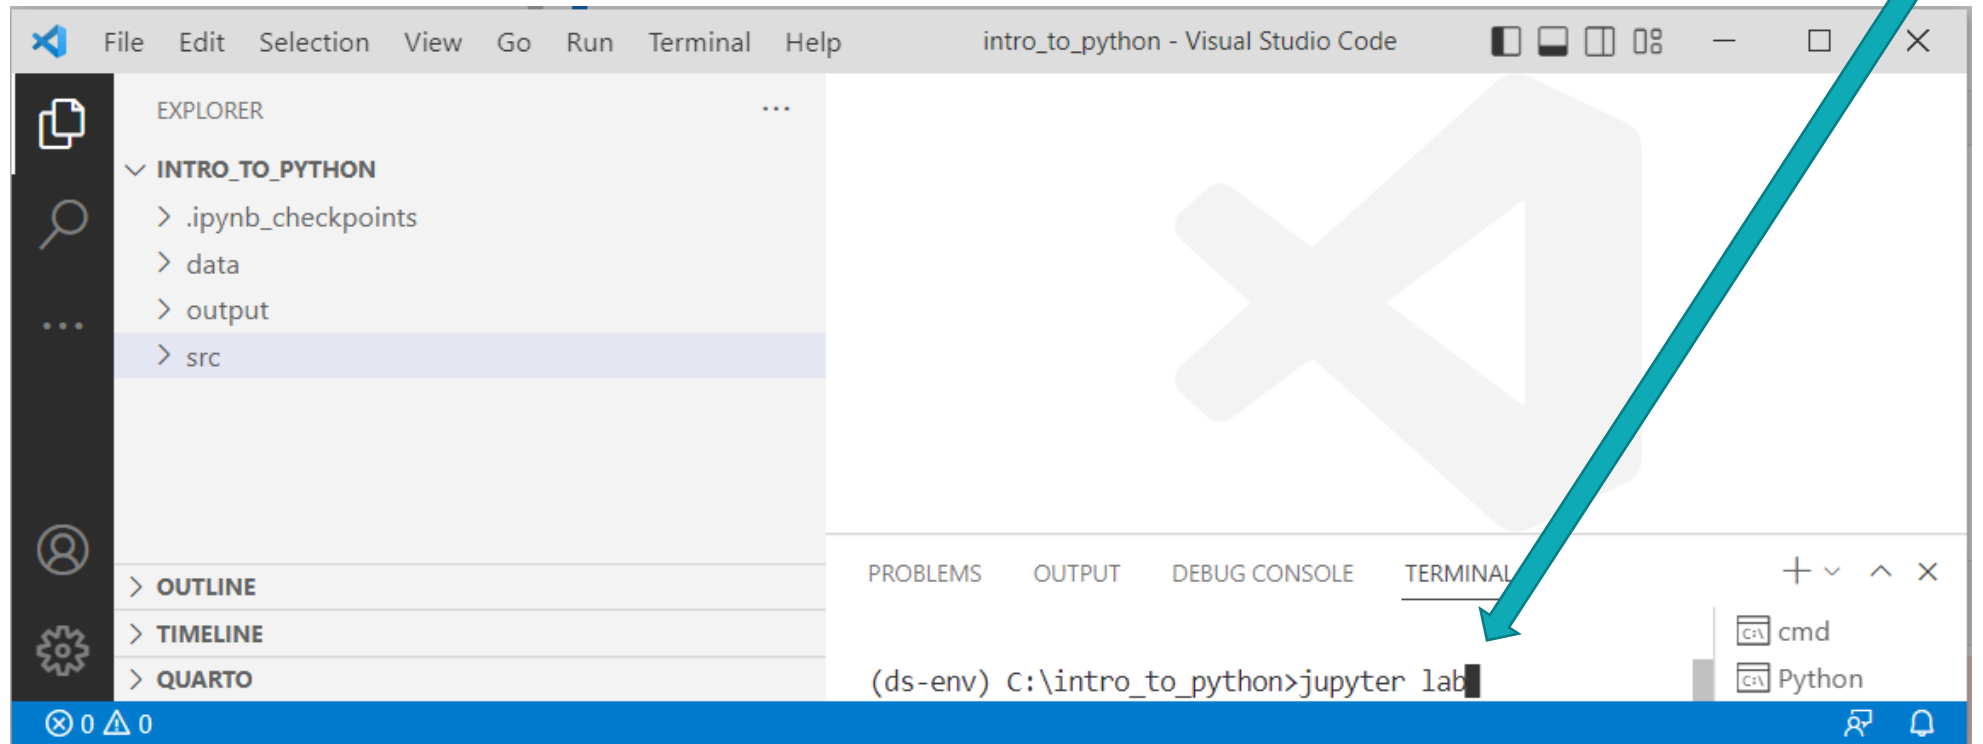

Challenge 2

If output is generally considered the most usable thing for end-users of data, why should it be considered as disposable?




Getting Help

Let's get coding!




Challenge 3

You want to find out what the **print()** function does. How would you look to answer this question? What does it do, does it send a print job to a printer?




Challenge 4

Someone asks you what the term “iterable” means in the context of a python function? How would you go about finding an answer? What does it mean?



Coding in Style

Let's get coding!



Challenge 5

A decorative swirl graphic on the left side of the slide, consisting of four circles connected by a line. The circles are colored purple, blue, light blue, and green from top to bottom. The number 1 is inside the purple circle.

1

Based on what we have covered, what, if any are the style code issues.

Example 1
`print (last_name)`

A decorative swirl graphic on the left side of the slide, consisting of four circles connected by a line. The circles are colored purple, blue, light blue, and green from top to bottom. The number 2 is inside the blue circle.

2

Example 2
`celcius = (fahrenheit - 32) / 1.8`

A decorative swirl graphic on the left side of the slide, consisting of four circles connected by a line. The circles are colored purple, blue, light blue, and green from top to bottom. The number 3 is inside the light blue circle.

3

Example 3
`index = ['snail', 'pig', 'elephant', 'rabbit', 'giraffe', 'coyote', 'horse', 'giraffe', 'ping pong']`

A decorative swirl graphic on the left side of the slide, consisting of four circles connected by a line. The circles are colored purple, blue, light blue, and green from top to bottom. The number 4 is inside the green circle.

4

CORE CONCEPTS



Challenge 1

So far in the notebook we should have:

```
variable_name = 'variable_value'  
first_name = 'Loki'  
age = 1054
```


In the next cell, run the following command:

```
print(last_name)
```

What happens?

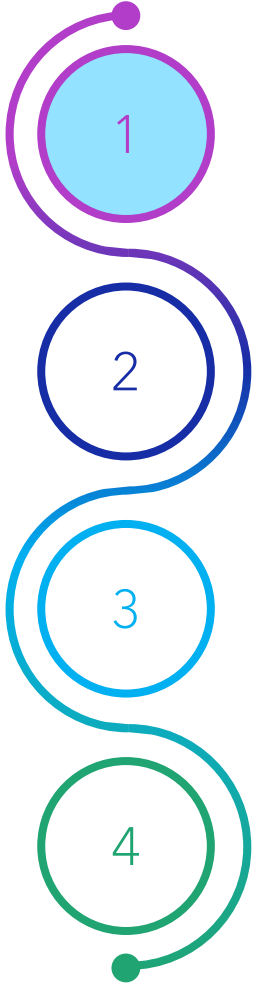
Why?

How can we fix it?



Challenge 2

Fill in the table below, showing the values of each variable in this program *after* each statement is executed.



Command	Value of x	Value of y	Value of swap
x = 1.0	1.0	Undefined	Undefined
y = 3.0	1.0	3.0	Undefined
swap = x	1.0	3.0	1.0
x = y	3.0	3.0	1.0
y = swap	3.0	1.0	1.0


Challenge 3

Given the following string:

```
full_name = 'Peregrin Fool of a Took Pippin'
```

What would these expressions return?

1. `full_name[2:8]`
2. `full_name[11:]` (without a value after the colon)
3. `full_name[:4]` (without a value before the colon)
4. `full_name[:]` (just a colon)
5. `full_name[11:-3]`
6. `full_name[-5:-3]`
7. What happens when you choose a stop value which is out of range?
(i.e. try `full_name[0:42]` or `full_name[:103]`)



Lists vs Dictionaries

Lists: use an ordered index, starting at 0

0	1	2	3	4
apple	banana	orange	potato	otter
A	'key'	'color'	99	'cute'

Dictionary: uses an unordered index, using values supplied by you, the coder

list[2] → orange

Error! list['color']

dictionary[2] Error!

dictionary['color'] → orange

Data Types



Integers (int)

positive or negative whole numbers

42, -7, 90_210, 0, ...

Floats (float)

real fractional numbers

3.14159, -87.6, 0.0, ...

Strings (str)

text written either in 'single' or "double" quotes

'Hello, World!', "42", 'zero'

Booleans (bool)


the logical values of True and False

Challenge 4

Use the `help()` and `round()` functions to print out the value of:

2.71828182845904523536

to 0 and 2 decimal places



Challenge 5

What will be the output of the following block of code?

```
num = 42  
animal = 'otter'
```

```
if num==42 and animal=='mouse':  
    print('Correct, that is the animal that found the answer')
```


```
elif num==42 and animal!='mouse':  
    print('Almost, the number is correct but not the animal.')  
    animal = 'dolphin'
```

```
elif num!=42 or animal=='mouse':  
    print('Almost, the animal is correct but not the number.')  
    num = 5
```

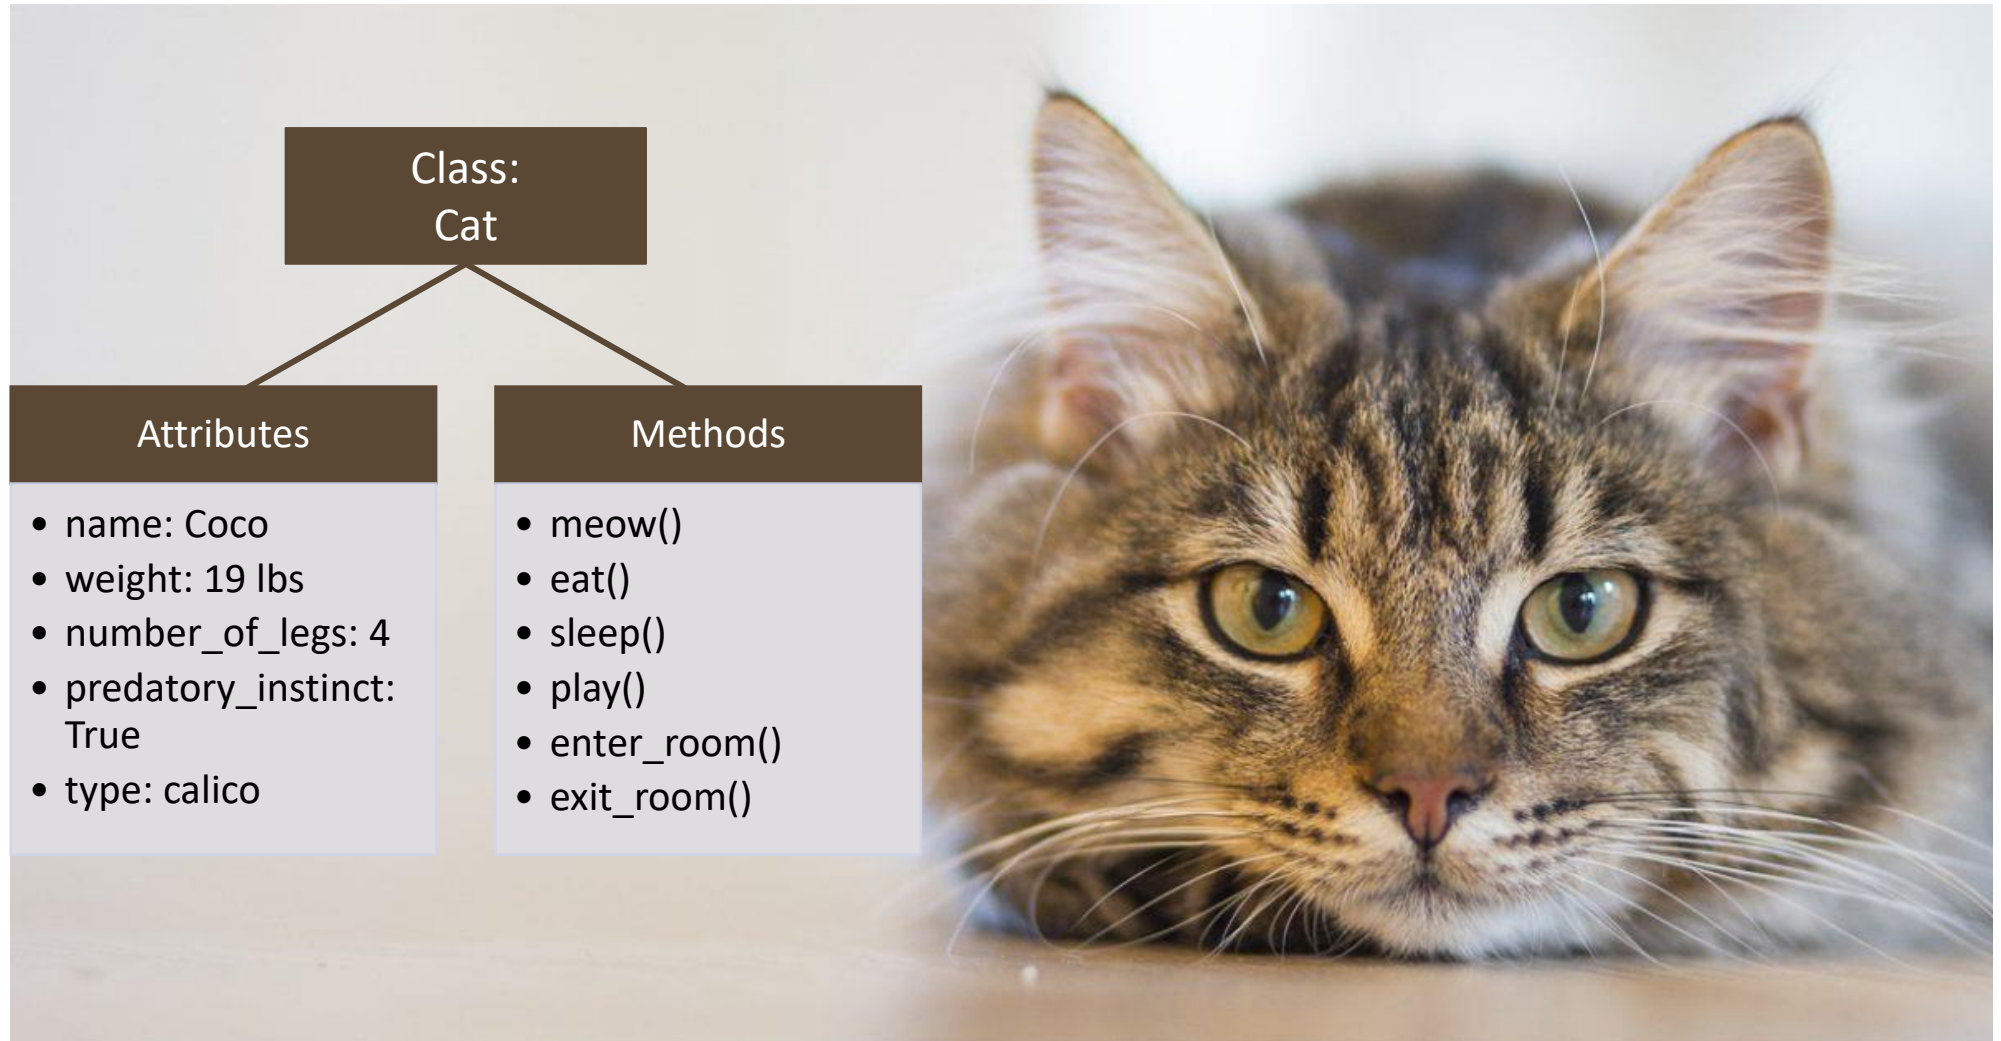

```
elif (1>3) or (4>3):  
    print('This has nothing to do with it, we just needed an or statement')
```

```
else:  
    print('Not even close, those pesky mice need to work harder.')  
    num = 19  
    animal = 'kangaroo'
```

```
print('The end result is the number', num, 'and animal', animal)
```



Objects in Python




Challenge 6

We currently have the string:

```
my_string = 'Peter Piper Picked a Peck of Pickled Peppers'
```

What happens if we try to run the following block of code:

```
print(my_string.isnumeric().upper())
```



Challenge 7


A common data string you see across government are various personal IDs. One example is the Personal Education Number (PEN), which is a nine-digit number assigned to you when you enter the K-12 School System. Oftentimes when looking at such an ID, any leading zeros that are an actual part of the PEN get stripped away, leading to unevenly sized strings of IDs across the dataset.

Try your code out with the following PENs:

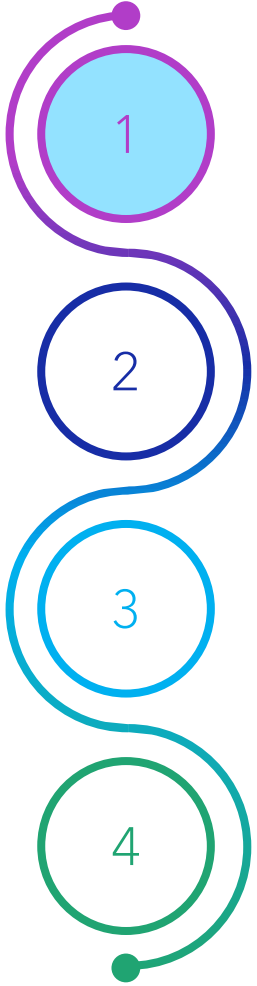
```
pen_1 = '12345678x'  
pen_2 = '123456789'  
pen_3 = '1234567890'  
pen_4 = '123456789'  
pen_5 = '123456'
```

Write a piece of code for PEN IDs that does the following:

- Checks to make sure that the ID is entirely numeric. If it is not, print out a warning that this is an invalid PEN.
- If the PEN is numeric, make sure that it is less than or equal to 9 digits long. If it is longer, print out a warning that this is an invalid PEN.
- If the PEN is too short, make sure to pad it with the appropriate number of 0's to the left. Print out the correct PEN.



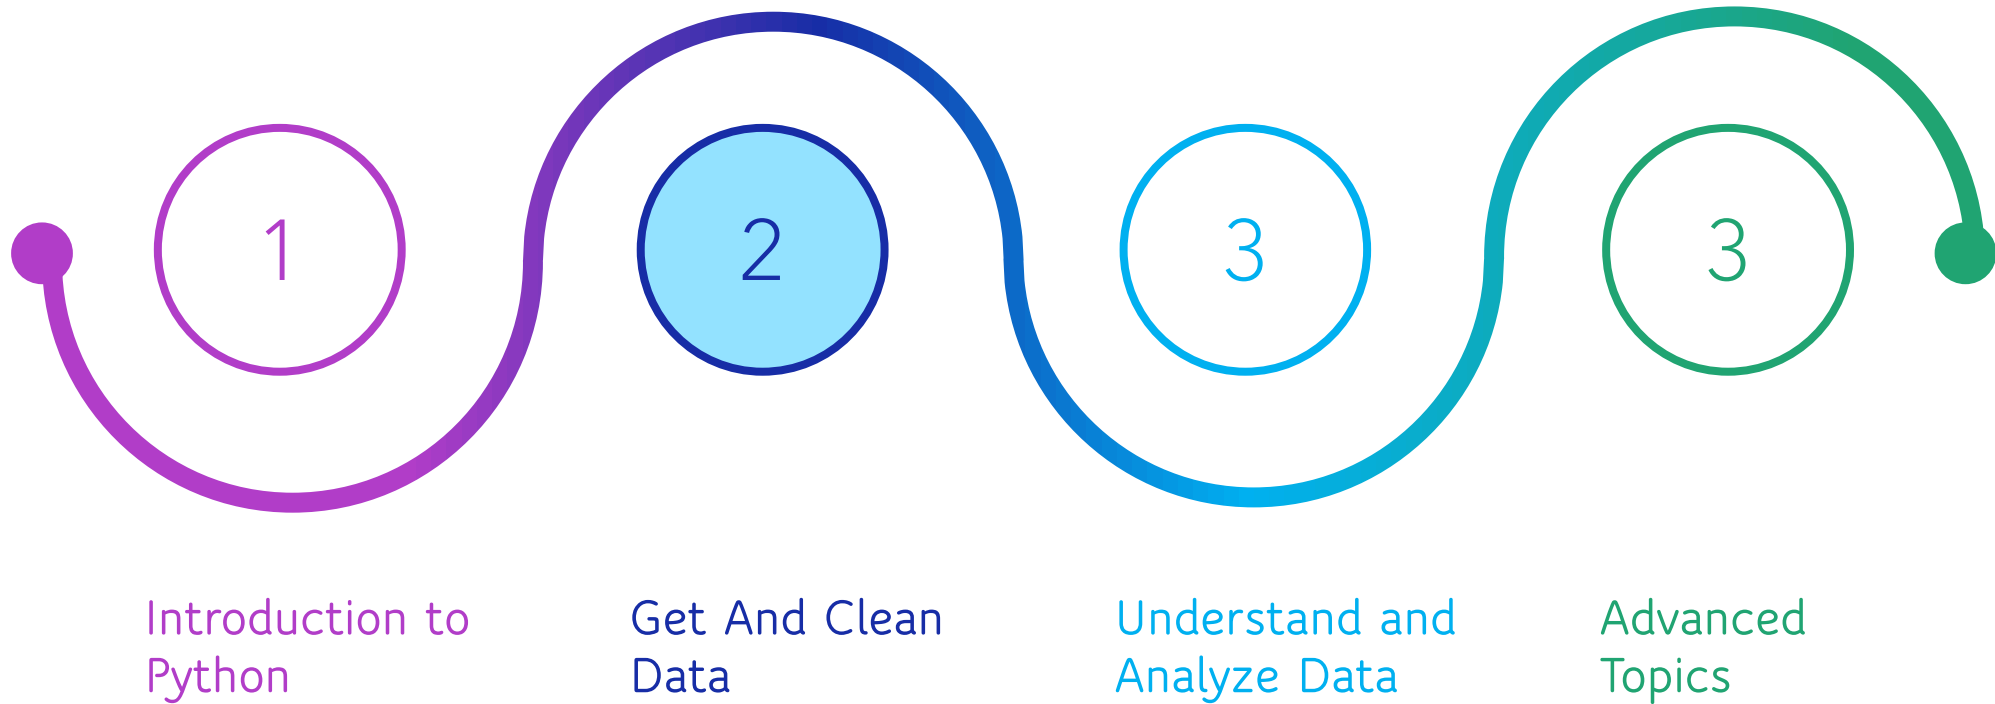
Common Package Aliases:



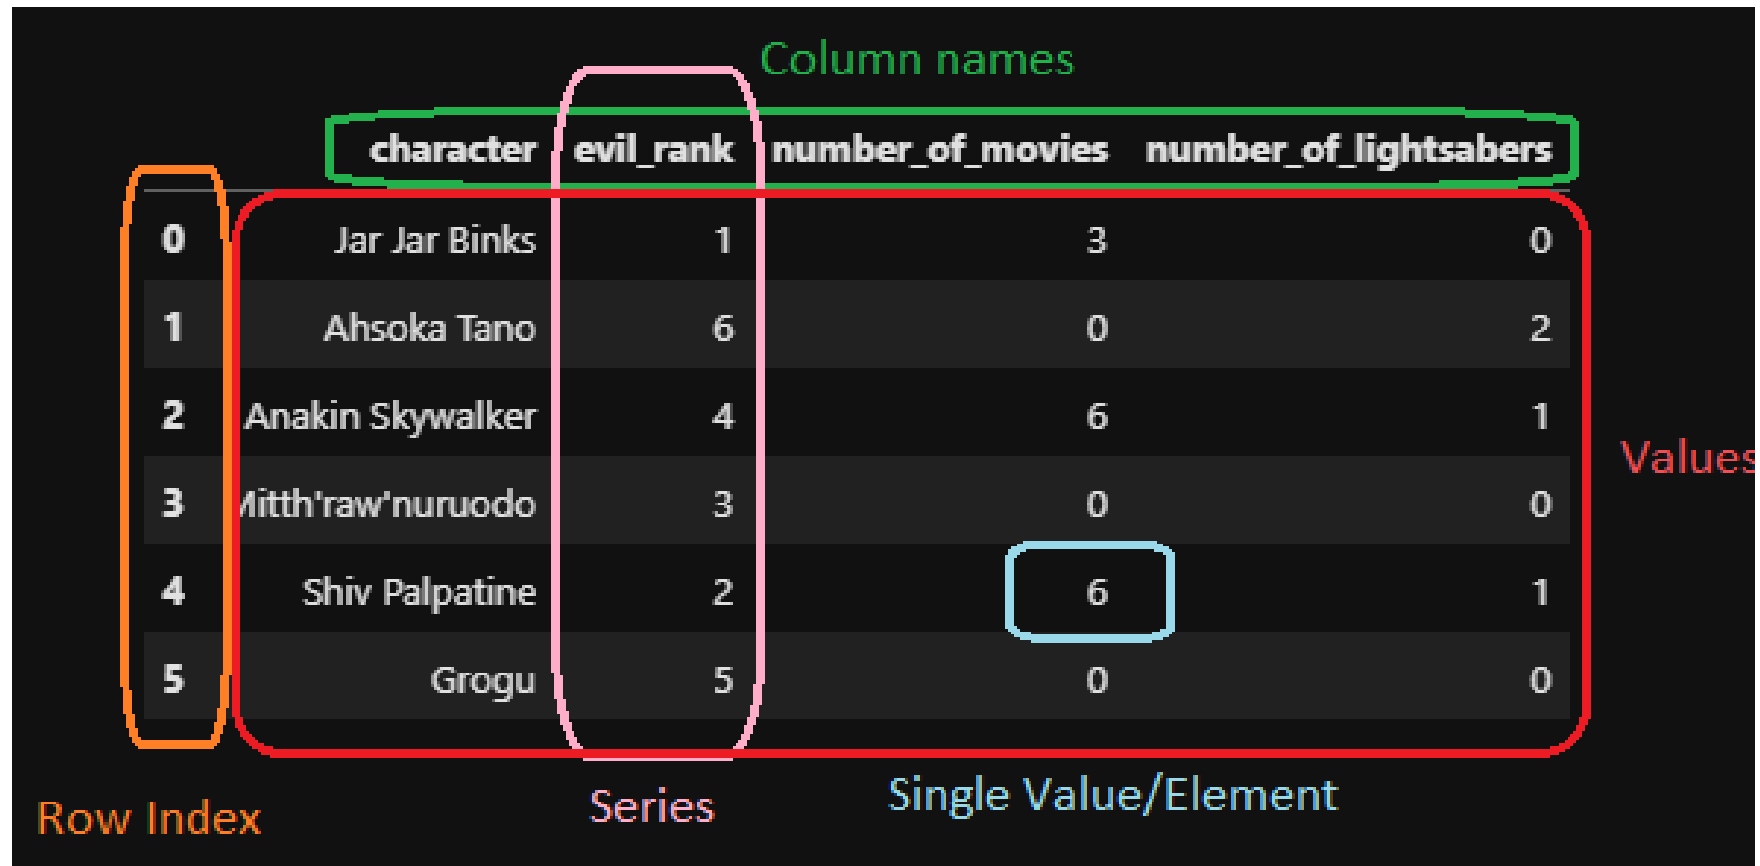
Package	Alias
pandas	pd
matplotlib.pyplot	plt
seaborn	sns
numpy	np

import PACKAGE **as** ALIAS

GETTING DATA WITH PANDAS



Pandas Dataframes



The diagram illustrates the components of a Pandas DataFrame using a table of Star Wars characters. The table is annotated with labels and boxes to identify its parts:

- Column names:** Indicated by a green box at the top of the table.
- Row Index:** Indicated by an orange box on the left side of the table.
- Series:** Indicated by a pink box encompassing the data columns.
- Single Value/Element:** Indicated by a light blue box around the value '6' in the 'number_of_movies' column for the row with index 4.
- Values:** Indicated by a red box encompassing the entire data area of the table.

	character	evil_rank	number_of_movies	number_of_lightsabers
0	Jar Jar Binks	1	3	0
1	Ahsoka Tano	6	0	2
2	Anakin Skywalker	4	6	1
3	Mitth'raw'nuruodo	3	0	0
4	Shiv Palpatine	2	6	1
5	Grogu	5	0	0

Challenge 1

Bring the gapfinder dataset back into pandas.

This time:

- Use the country as the index
- Only include the year, continent and population columns
- **Hint: ?pd.read_csv**
- **Hint 2: What do the arguments `index_col` and `usecols` do?**

Challenge 2


We have the gapfinder dataset loaded into an object called **df**

What will happen if we run the following code:

```
df.sample(42).describe()
```

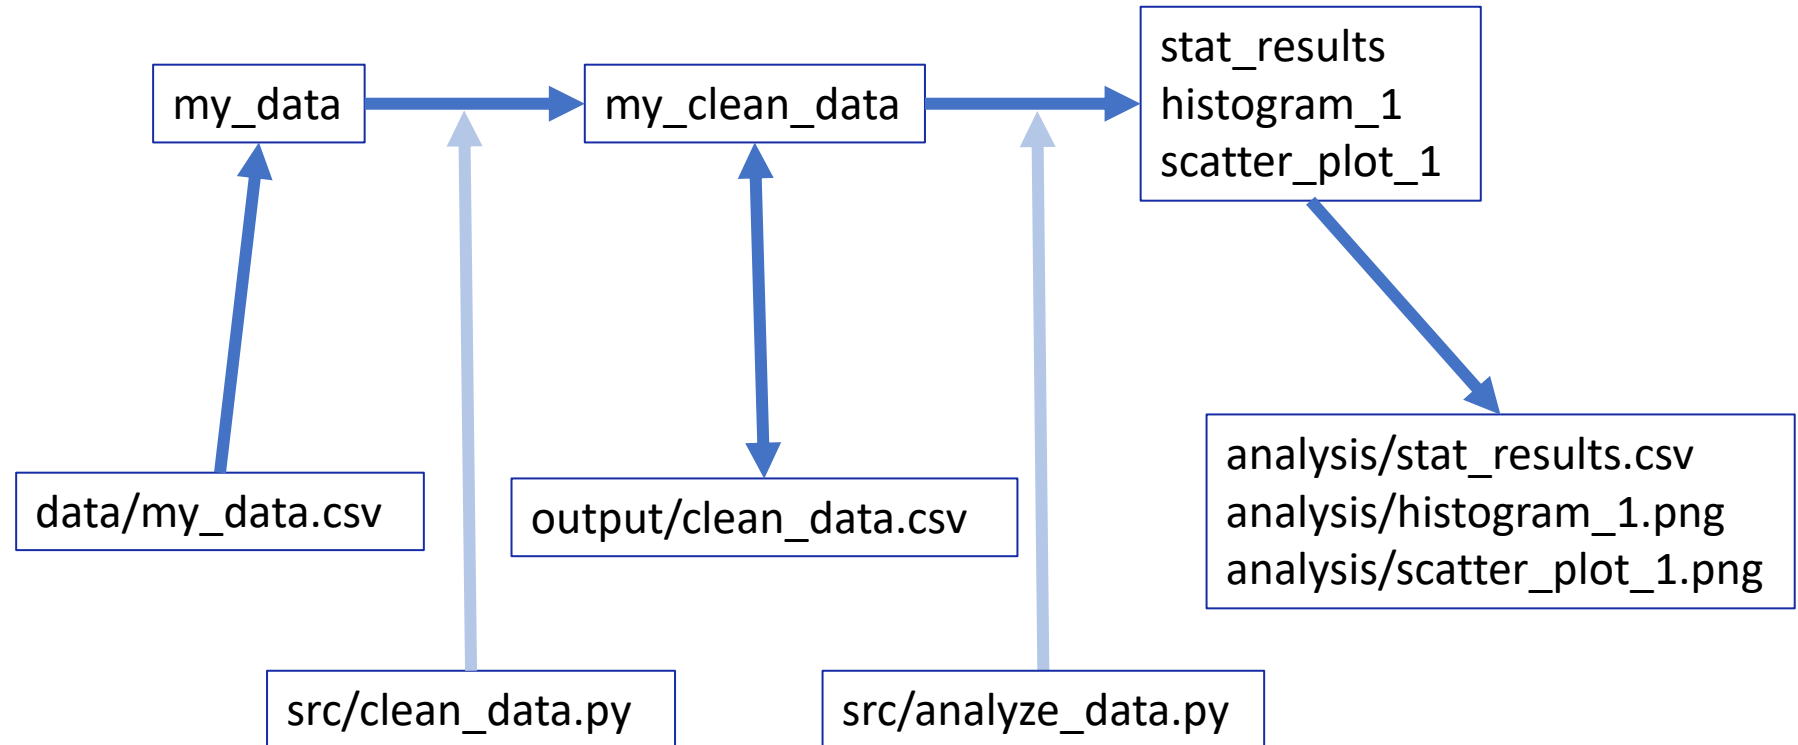
1. Do we expect the results to be the same as **df.describe()**? Why or why not?
2. Run the code again. Are the results the same or different than before? Can you explain?

Beginning of a Reproducible Pipeline




Python

**Project
Folder**



Challenge 3



1

Save a copy of all the summary statistics for the gapfinder dataset.

2

Only include the statistics for the **pop** and **lifeExp** columns.

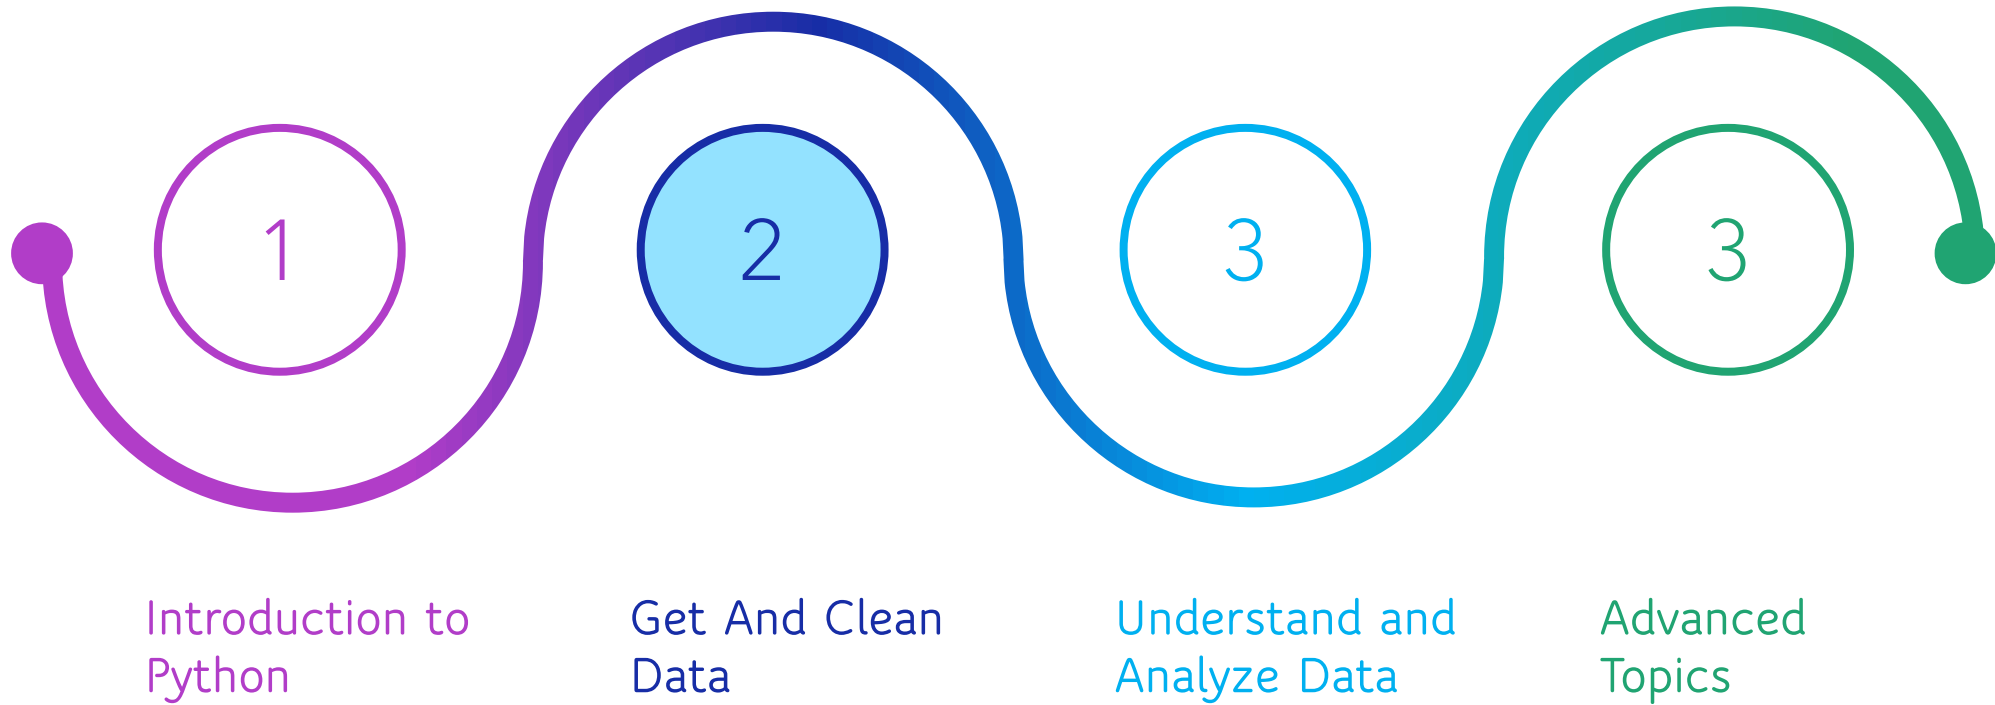
3

What happens when we include or exclude the index?

4

Hint: columns, index

CLEANING DATA



Challenge 1

You are given a dataset for several days time frame in Feb/March 2021 of the average daily temperature in celcius for three Canadian cities. You would like to ultimately create a plot of the temperatures but don't want to have gaps in the data. Run the code below to get the data and see what it looks like.


```
import pandas as pd
```

```
temps = "https://raw.githubusercontent.com/bcgov/" \
        "ds-intro-to-python/main/data/citytemps.csv"
```


```
city_temps = pd.read_csv(temps)
```

```
city_temps.head(10)
```

Unfortunately you notice that several of the days have data missing. Write code that leverages the `interpolate()` function to fill in the missing data. Run the code. Is this an appropriate solution for the problem? What are the pros and cons to this approach.



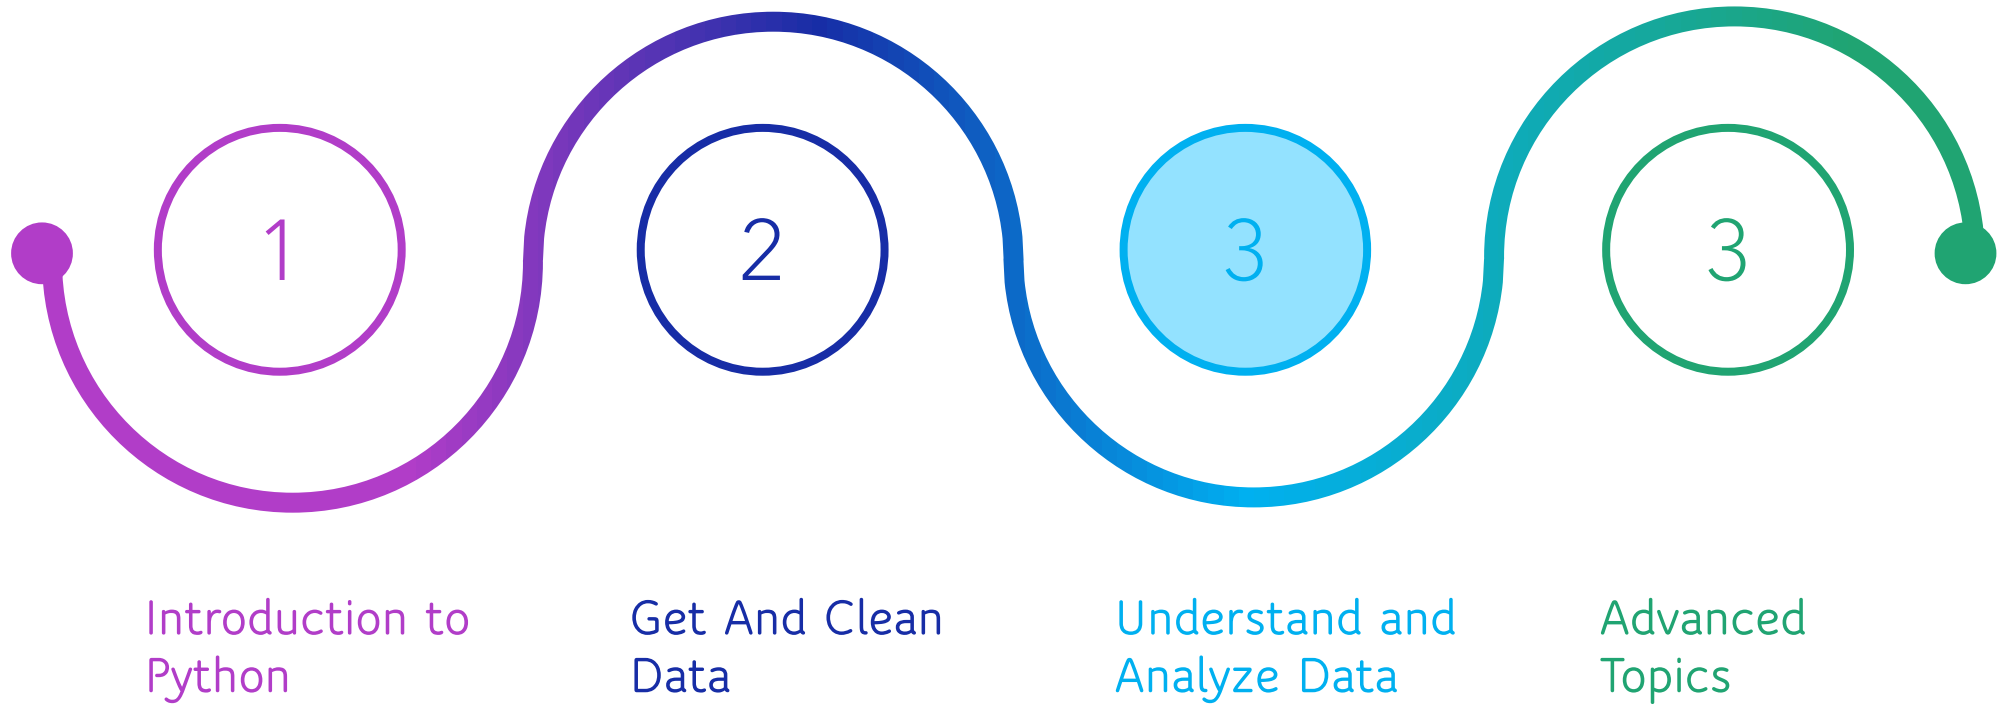
Challenge 2



You would like to clean up the scale used in the question about whether one's mental illness affects one's work.


- Find the column (pandas series) that needs to be transformed and transform its values using `replace()` with numerical values that make sense.
- Display the modified dataframe showing the replaced values.

EXPLORING DATA STRUCTURES




Challenge 1

Let's say you want to narrow down the dataset to include only the country and the year. How would you do it? Don't forget to run your code so it also shows a view of the result so you can confirm the code worked as you wanted it to.




Challenge 2

Your director has come to you and asked if you know what the life expectancy has been in Canada since 1992. How would you use pandas code to get the data you need? And after having run the code, what's the answer?




Challenge 3

Your director has come back to you and wondered about whether the life expectancy of people changed during the 1970s in Cambodia. Use what you know about selecting rows and sorting data to get the data you need to answer the question, sorted from most recent at the top.

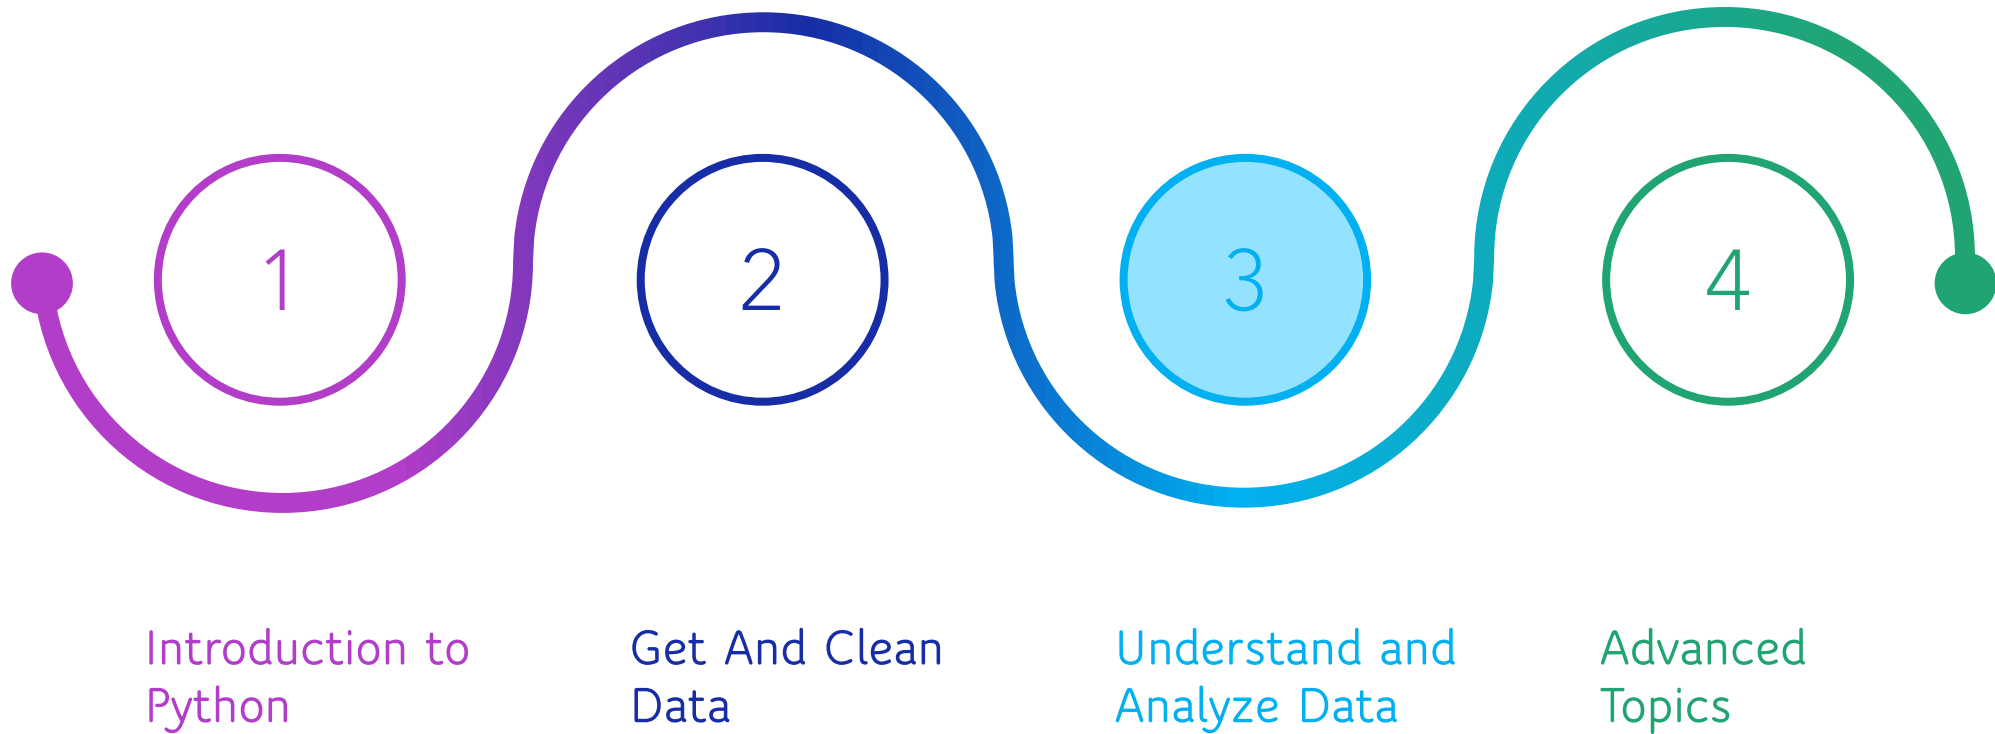


Challenge 4

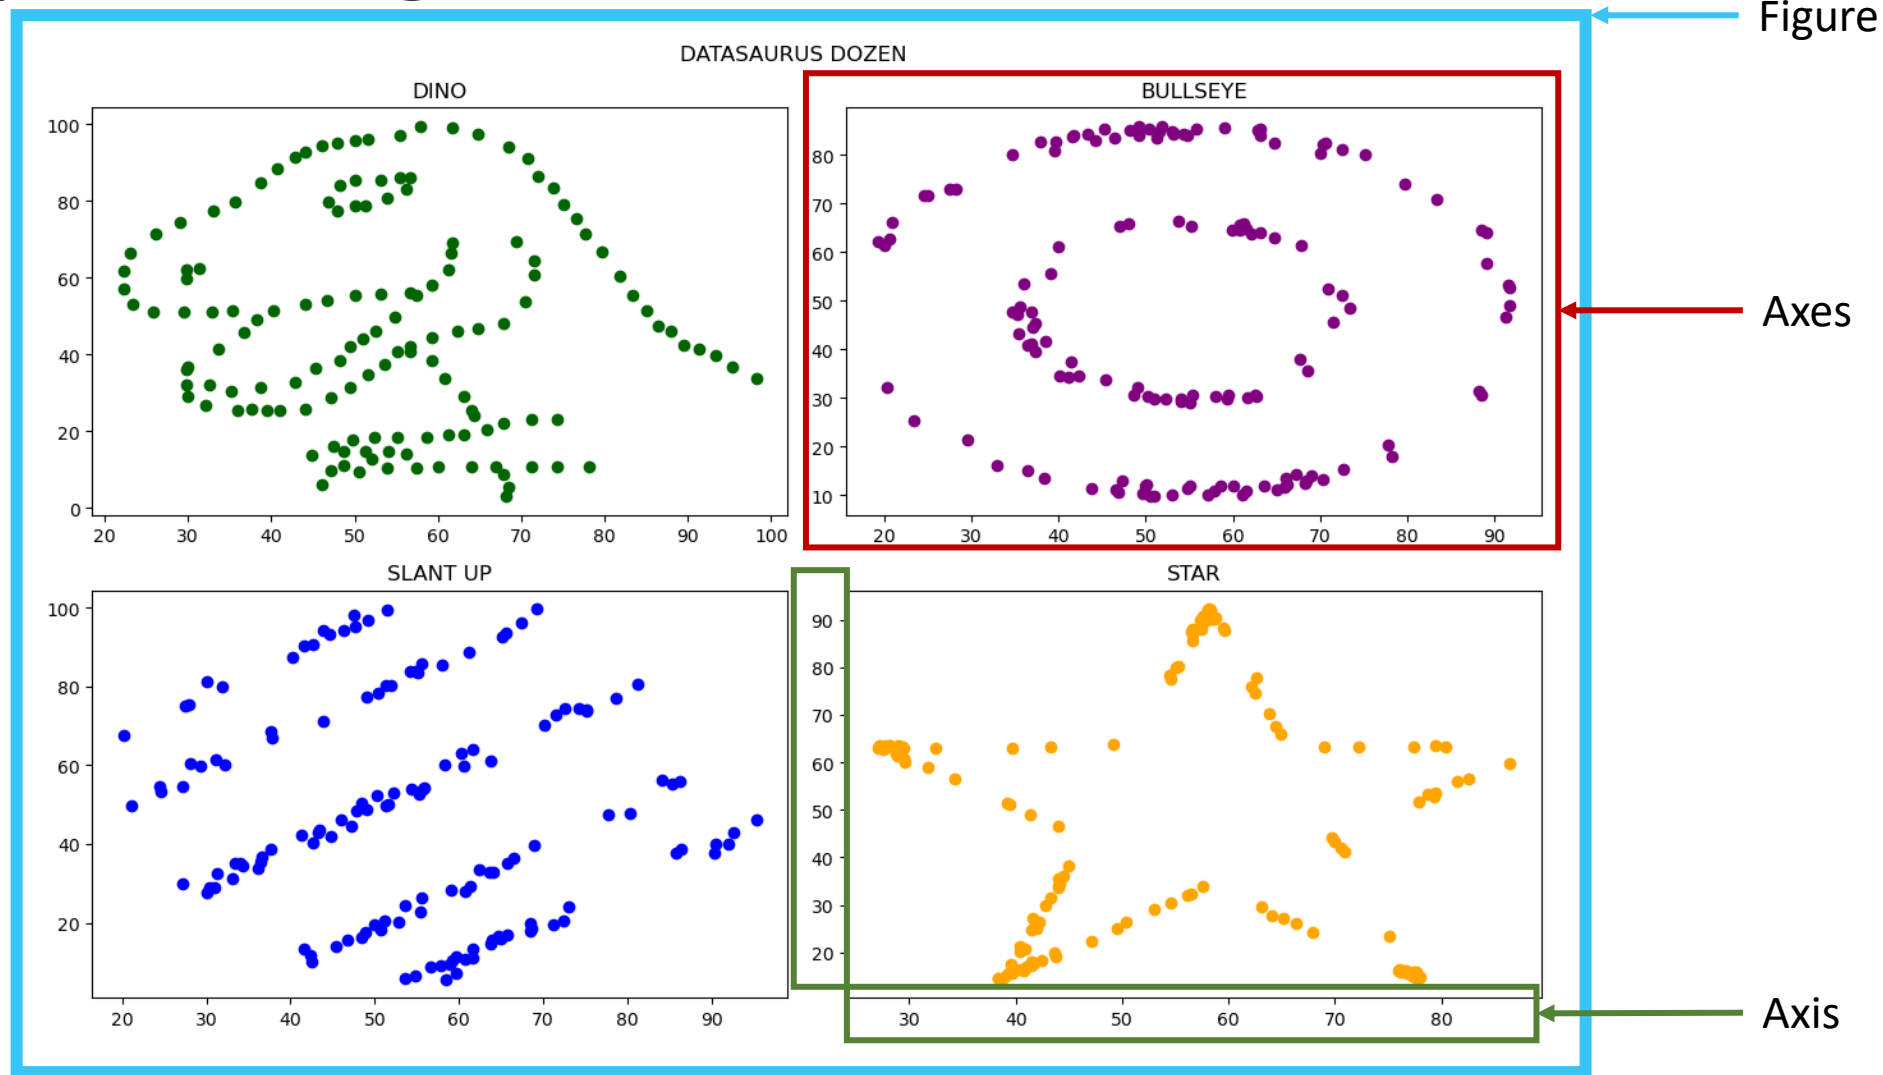

You would like to summarize population as well as life expectancy by year, grouped by continent. Pick some aggregations that would make sense to look at for this task. Don't worry about re-setting the index or about creating new labels for the result.



GRAPHICAL DEPICTIONS OF DATA




Matplotlib Figures



Matplotlib Figures

1. Entire figure and axes is initialized when we create the first plot. These are assigned to the **current figure/axes**.
2. Add secondary/tertiary plots to the same axes line by line.
 - Individual plot characteristics are defined on each of these lines (label, plot type, color used, etc.)
3. Label each **axes** as required
 - Titles, legend names, legend values, etc.
4. Save the plot
5. Display the plot



Challenge 1



1

Using what we have learned so far about pandas filtering and matplotlib functionality, produce a plot of bill lengths vs. depths for the penguins dataset.

2

Colour the three penguin species (Adelie, Chinstrap, Gentoo) different colours.

3

Include a legend that identifies which species is which colour!

4

HINT: use what we learned with Stu to filter the dataframe to a single species (create three new dataframes...)

Seaborn Plots

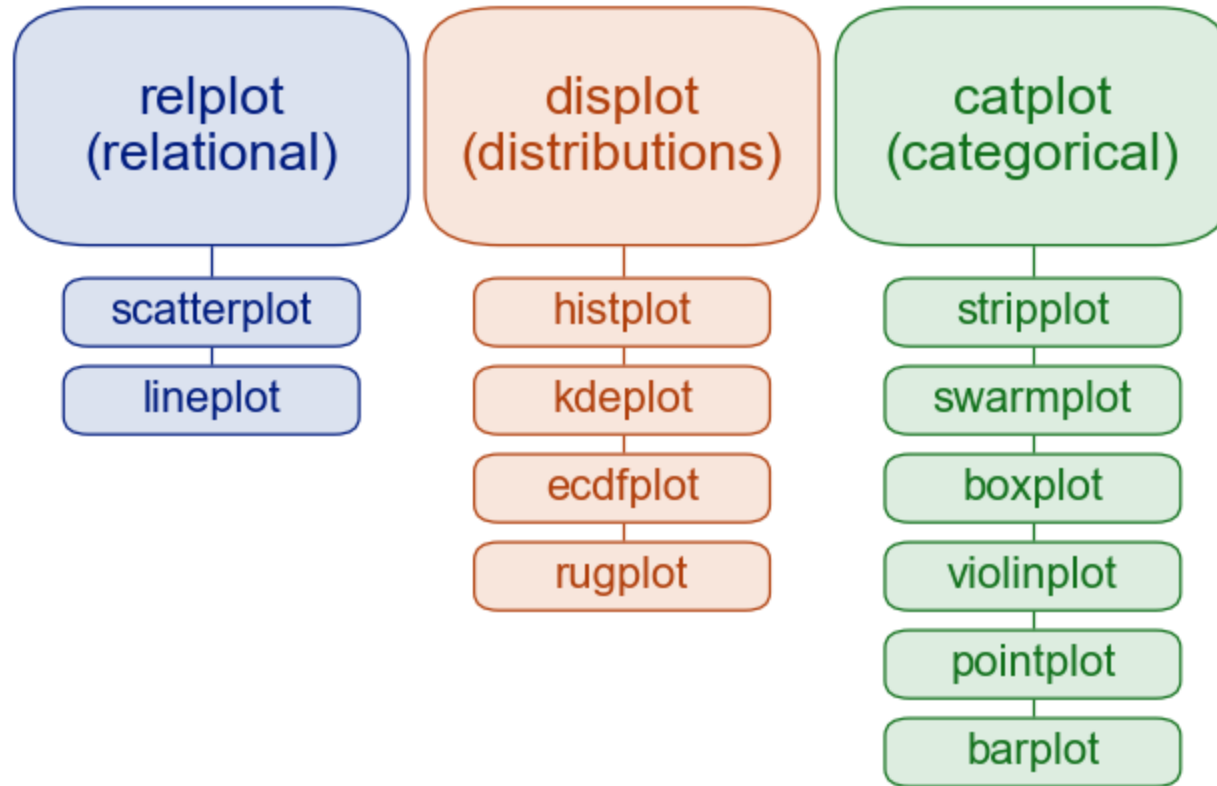




Figure Level:

- Simple, single command for multiple different plot types
- Returns a 'fully formed' final figure
- Cannot easily access the underlying matplotlib components

Axes Level:

- Individualized for specific plot types
- Returns a single axes that can be added to other matplotlib figures
- Easier to customize, but more 'fidgety'

Relational Plots



Usage

- Let us look for patterns in a dataset between two or more numerical features
- Can be used to identify trends over time
- Can be used to display uncertainty in a model

Types

- Scatter plots
- Line plots

Conditional Options

- style, hue, size
- row, col


Challenge 2



1 Using the Seaborn **load_dataset** method, import a dataset called 'dowjones'.

- 2 1. What does this dataset include/represent?
- 3 2. On what date does the data represented here reach its highest price?
- 4 3. Create a line plot to depict the data. Does the plot agree with the answer you came up with in part 2?

Challenge 3





1 Choose any two numerical features in the penguins dataset.

2 Produce a scatter plot that highlights the differences between species.

3 Separate the results into two plots – one for male and one for female penguins.

4

Distribution Plots



Usage

- Analyze the data to see how features are distributed
- Identify outliers
- Compare subsets

Types

- Histograms
- KDE (Kernel Density Estimation)
- ECDF (Empirical Cumulative Distribution Function)

Conditional Options

- hue
- row/col

Challenge 4



1

Create a KDE plot of body mass from the penguins dataset. (body_mass_g)

2


Does it appear to be bi-modal (2 peaks)?

3

If so, create another plot (or more) to identify what may be the cause. Are male penguins heavier? Is it a specific species? Does the island matter?

4

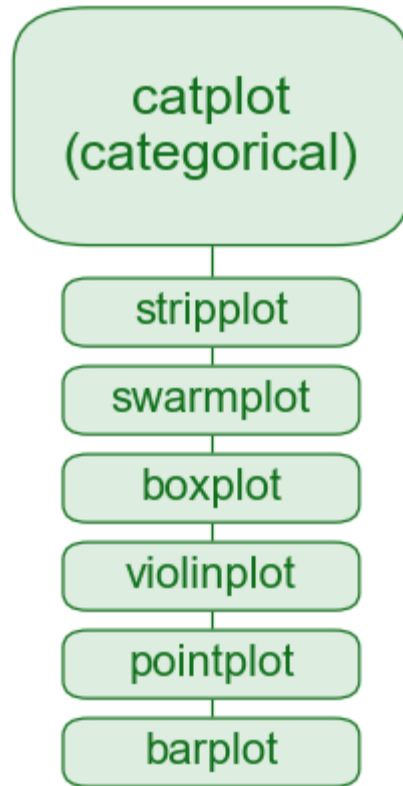

Challenge 5



1 Which penguin species was most frequently included in this dataset?

2 Display your result visually!

Categorical Plots



Usage

- Let us look for patterns in a dataset using at least one categorical feature
- Default is to display **aggregations** of the data for most types!

Types

- Categorical scatterplots: swarm, strip
- Categorical distribution plots: box, violin, boxen
- Categorical estimate plots: point, bar, count

Conditional Options

- hue
- row/col